


Lecture 19: Multistep decisions


Instructor: Sergei V. Kalinin

Automated Experiment At UT Knoxville


Automated SPM



High-throughput synthesis




Automated Electron Microscopy



Gerd Duscher

Robotic film deposition/characterization



Mahshid Ahmadi

Multistage decisions

There is no computation without representation!

- Learn algorithms: from simple to complex ones
- Learn new concepts that emerge as a part of ML, e.g. VAE, decision making, etc.
- Learn how to represent the system in the ML setting
- Learn how to formulate the reward functions

Multistage decisions: agent

Autonomous car (or robot, or even Roomba):

- **Percepts:** inputs from sensors, cameras, GPS system, etc
- **Actions:** steering, acceleration, controls
- **Goals:** reach destination, safety, comfort, maximize profit, etc...
- **Environment:** street, highway, mountain trail,
- Goals specifiable by performance measure defining a numerical value for any environment history (reward)
- **Rational action:** whichever action maximizes the expected value of the performance measure given the percept sequence to date
- **Rational** does not mean omniscient, clairvoyant, or successful
- Humans are (often) not rational (stock market, tests)

Multistage decisions

	Solitaire	Backgammon	Internet shopping	Taxi
<u>Accessible</u> ??	Yes	Yes	No	No
<u>Deterministic</u> ??	Yes	No	Partly	No
<u>Episodic</u> ??	No	No	No	No
<u>Static</u> ??	Yes	Semi	Semi	No
<u>Discrete</u> ??	Yes	Yes	Yes	No

- The environment type determines the agent design
- Real world is inaccessible, stochastic, sequential, dynamic, continuous


Agent functions

- Agent function is completely specified by policy (agent function) mapping percept sequence to actions
- Most primitive policy will be look-up table (i.e. simple enumeration)


Agent types:

- Simple reflex agents
- Reflex agents with state
- Goal-based agents
- Utility-based agents


Simple reflex agents




Reflex agents with state



Goal based agents



Utility based agents



From agents to search

- Rational agents need to perform sequences of actions to achieve goals
- Intelligent behavior can be generated by having a look-up table or reactive policy that tells the agent what to do in every circumstance, but:
 - Such a table or policy is difficult to build
 - All contingencies must be anticipated
- A more general approach is for the agent to have knowledge of the world and how its actions affect it and be able to simulate execution of actions in an internal model of the world in order to determine a sequence of actions that will accomplish its goals
- This is the general task of problem solving and is typically performed by searching through an internally modelled space of world states

From agents to search

- **Given:**
 - An initial state of the world
 - A set of possible actions or operators that can be performed.
 - A goal test that can be applied to a single state of the world to determine if it is a goal state.
- **Find:** A solution stated as a path of states and operators that shows how to transform the initial state into one that satisfies the goal test.
- The initial state and set of operators implicitly define a state space of states of the world and operator transitions between them. May be infinite.


From agents to search

- **Path cost:** a function that assigns a cost to a path, typically by summing the cost of the individual operators in the path. May want to find minimum cost solution.
- **Search cost:** The computational time and space (memory) required to find the solution.
- Generally there is a trade-off between path cost and search cost and one must satisfy and find the best solution in the time that is available

From agents to search

- Rational agents need to perform sequences of actions to achieve goals
- Intelligent behavior can be generated by having a look-up table or reactive policy that tells the agent what to do in every circumstance, but:
 - Such a table or policy is difficult to build
 - All contingencies must be anticipated
- A more general approach is for the agent to have knowledge of the world and how its actions affect it and be able to simulate execution of actions in an internal model of the world in order to determine a sequence of actions that will accomplish its goals
- This is the general task of problem solving and is typically performed by searching through an internally modelled space of world states

Route finding



- Initial state: Arad
- Goal state: Bucharest
- Path cost: Number of intermediate cities, distance traveled, expected travel time

From agents to search

- Constraints and constraint-satisfaction problems

- 8-puzzle (sliding tile puzzle)


5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

- 8-queens problem (N-queens problem)



More realistic problems:

- Route finding (Google Maps)
- Travelling salesman problem
- Supply chain planning
- Integrated circuit design
- Experiment planning


Search concepts

- A state can be expanded by generating all states that can be reached by applying a legal operator to the state
- State space can also be defined by a successor function that returns all states produced by applying a single legal operator
- A search tree is generated by generating search nodes by successively expanding states starting from the initial state as the root
- A search node in the tree can contain
 - Corresponding state
 - Parent node
 - Operator applied to reach this node
 - Length of path from root to node (depth)
 - Path cost of path from initial state to node


(a) The initial state

Arad

(b) After expanding Arad



(c) After expanding Sibiu



Search strategies

- Easiest way to implement various search strategies is to maintain a queue of unexpanded search nodes
- Different strategies result from different methods for inserting new nodes in the queue

Search strategies

Properties of search strategies


- Completeness: does it always find solution if one exists?
- Time Complexity: number of nodes generated/expanded
- Space Complexity: maximum number of nodes in memory
- Optimality: does it always find least cost solution?

Informed vs. uninformed:

- Uninformed search strategies (blind, exhaustive, brute force) do not guide the search with any additional information about the problem.
- Informed search strategies (heuristic, intelligent) use information about the problem (estimated distance from a state to the goal) to guide the search

Breadth-first search


- Expands search nodes level by level, all nodes at level d are expanded before expanding nodes at level $d+1$



- Implemented by adding new nodes to the end of the queue
- Since eventually visits every node to a given depth, guaranteed to be complete
- Optimal provided path cost is a nondecreasing function of the depth of the node (e.g. all operators of equal cost) since nodes explored in depth order

Breadth-first search

- Expands search nodes level by level, all nodes at level d are expanded before expanding nodes at level $d+1$



- Implemented by adding new nodes to the end of the queue
- Since eventually visits every node to a given depth, guaranteed to be complete
- Optimal provided path cost is a nondecreasing function of the depth of the node (e.g. all operators of equal cost) since nodes explored in depth order

Complexity of breadth-first search

- Assume there are an average of b successors to each node, called the branching factor

- Therefore, to find a solution path of length d must explore

$$1 + b + b^2 + b^3 + \dots + b^d$$


- Plus need b^d nodes in memory to store leaves in queue
- Assuming can expand and check 1000 nodes/sec and need 100 bytes/node storage, $b=10$

Depth	Nodes	Time	Memory
0	1	1 millisecond	100 bytes
2	111	.1 seconds	11 kilobytes
4	11,111	11 seconds	1 megabyte
6	10^6	18 minutes	111 megabytes
8	10^8	31 hours	11 gigabytes
10	10^{10}	128 days	1 terabyte
12	10^{12}	35 years	111 terabytes
14	10^{14}	3500 years	11,111 terabytes

Note memory is a bigger problem than time.

Uniform cost search


- Like breadth-first except always expand node of least cost instead of least depth (i.e. sort new queue by path cost)



- Do not recognize goal until it is the least cost node on the queue and removed for goal testing
- Therefore, guarantees optimality as long as path cost never decreases as a path increases (non-negative operator costs)

Depth-first search

- Always expand node at deepest level of the tree, i.e. one of the most recently generated nodes. When hit a dead-end, backtrack to last choice



Breadth-first search

- Not guaranteed to be complete since might get lost following infinite path
- Not guaranteed optimal since can find deeper solution before shallower ones explored
- Time complexity in worst case is still $O(b^d)$ since need to explore entire tree. But if many solutions exist may find one quickly before exploring all of the space
- Space complexity is only $O(b^m)$ where m is maximum depth of the tree since queue just contains a single path from the root to a leaf node along with remaining sibling nodes for each node along the path
- Can impose a depth limit, l , to prevent exploring nodes beyond a given depth. Prevents infinite regress, but incomplete if no solution within depth limit.

Iterative deepening

- Conduct a series of depth-limited searches, increasing depth-limit each time.
- Seems wasteful since work is repeated, but most work is at the leaves at each iteration and is not repeated.

Depth-first:

$$1 + b + b^2 + \dots + b^{d-2} + b^{d-1} + b^d$$

Iterative deepening:


$$(d+1)1 + db + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

Heuristic Search


- Heuristic or informed search exploits additional knowledge about the problem that helps direct search to more promising paths.
- A **heuristic function**, $h(n)$, provides an estimate of the cost of the path from a given node to the closest goal state
- Must be zero if node represents a goal state.
 - Example: Straight-line distance from current location to the goal location in a road navigation problem
- Many search problems are NP-complete so in the worst case still have exponential time complexity; however a good heuristic can:
 - Find a solution for an average problem efficiently.
 - Find a reasonably good but not optimal solution efficiently.

Best first search

- At each step, best-first search sorts the queue according to a heuristic function



Best first search



- Does not find shortest path to goal (through Rimnicu) since it is only focused on the cost remaining rather than the total cost


Best first search properties

- Not complete, may follow infinite path if heuristic rates each state on such a path as the best option. Most reasonable heuristics will not cause this problem however
- Worst case time complexity is still $O(b^m)$ where m is the maximum depth.
- Since must maintain a queue of all unexpanded states, space-complexity is also $O(b^m)$
- However, a good heuristic will avoid this worst-case behavior for most problems.

Beam search


- Space and time complexity of storing and sorting the complete queue can be too inefficient
- Beam search trims queue to the best n options (n is called the beam width) at each point
- Focuses search more but may eliminate solution even for finite search graphs

Example for $n=2$



Hill climbing

- Beam search with a beam-width of 1 is called hill climbing
- Pursues locally best option at each point, i.e. the best successor to the current best node.
- Subject to local maxima, plateaux, and ridges.



Minimizing total cost: A* Search

- A* combines features of uniform cost search (complete, optimal, inefficient) with best-first (incomplete, non-optimal, efficient).
- Sort queue by estimated total cost of the completion of a path:

$$f(n) = g(n) + h(n)$$


Cost to this moment

Estimated cost to the goal

- If the heuristic function always underestimates the distance to the goal, it is said to be admissible.
- If h is admissible, then $f(n)$ never overestimates the actual cost of the best solution through n

A* Search Example

AIMA Slides ©Stuart Russell and Peter Norvig, 1998




- Finds the optimal path to Bucharest through Rimnicu and Pitesti

Optimality of A^*

If h is admissible, A^* will always find a least cost path to the goal

- Proof by contradiction:
 - Let G be an optimal goal state with a path cost f^*
 - Let G_2 be a suboptimal goal state supposedly found by A^*
 - Let n be a current leaf node on an optimal path to G




- Since h is admissible: $f^* \geq f(n)$
- If G_2 is chosen for expansion over n then: $f(n) \geq f(G_2)$
- Therefore: $f^* \geq f(G_2)$
- Since G_2 is a goal state, $h(G_2)=0$, therefore $f(G_2) = g(G_2)$ and $f^* \geq g(G_2)$
- Therefore G_2 is optimal. Contradiction.

Lemma: A^* expands nodes in order of increasing f value

Gradually adds “ f -contours” of nodes (cf. breadth-first adds layers)

Contour i has all nodes with $f = f_i$, where $f_i < f_{i+1}$



Other properties of A^*

- A^* is complete as long as
 - Branching factor is always finite
 - Every operator adds cost at least $\delta > 0$
- Time and space complexity still $O(b^m)$ in the worst case, since the search must maintain and sort complete queue of unexplored options.
- However, with a good heuristic can find optimal solutions for many problems in reasonable time
- Again, space complexity is a worse problem than time.

But what about heuristic functions?

- 8-puzzle search space
 - Typical solution length: 20 steps
 - Average branching factor: 3
 - Exhaustive search: $3^{20} = 3.5 \times 10^9$
 - Bound on unique states:
 $9! = 362,880$

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

- Admissible Heuristics:
 - Number of tiles out of place (h_1): 7
 - City-block (Manhattan) distance (h_2): $2+3+3+2+4+2+0+2=18$

Using heuristics:

- Experiments on sample problems can determine the number of nodes searched and CPU time for different strategies
- One other useful measure is effective branching factor: If a method expands N nodes to find solution of depth d , and a uniform tree of depth d would require a branching factor of b^* to contain N nodes, the effective branching factor is b^* , so that $N = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$

Experimental Results on 8-puzzle problems

d	Search Cost			Effective Branching Factor		
	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	364404	227	73	2.78	1.42	1.24
14	3473941	539	113	2.83	1.44	1.23
16	–	1301	211	–	1.45	1.25
18	–	3056	363	–	1.46	1.26
20	–	7276	676	–	1.47	1.27
22	–	18094	1219	–	1.48	1.28
24	–	39135	1641	–	1.48	1.26

Quality of heuristics

- Since A^* expands all nodes whose f value is less than that of an optimal solution, it is always better to use a heuristic with a higher value as long as it does not over-estimate.
- Therefore h_2 is uniformly better than h_1 or h_2 dominates h_1
- A heuristic should also be easy to compute, otherwise the overhead of computing the heuristic could outweigh the time saved by reducing search (e.g. using full breadth-first search to estimate distance wouldn't help).

Inventing Heuristics

- Many good heuristics can be invented by considering relaxed versions of the problem (abstractions)
- **For 8-puzzle:** A tile can move from square A to B if A is adjacent to B and B is blank
 - (a) A tile can move from square A to B if A is adjacent to B
 - (b) A tile can move from square A to B if B is blank
 - (c) A tile can move from square A to B
- If there are a number of features that indicate a promising or unpromising state, a weighted sum of these features can be useful. Learning methods can be used to set weights.